



Data Eng

```
In [1]: import findspark
findspark.init()

from pyspark.sql import SparkSession
from pyspark.sql import functions as F
from pyspark.sql.types import *
import pandas as pd

spark = SparkSession.builder \
    .appName("DE_DemandForecasting") \
    .master("local[*]") \
    .config("spark.sql.adaptive.enabled", "true") \
    .config("spark.sql.adaptive.coalescePartitions.enabled", "true") \
    .getOrCreate()

print("Spark Version:", spark.version)
```

```
26/05/02 00:43:16 WARN Utils: Your hostname, DESKTOP-8HUS7B0 resolves to a loop
back address: 127.0.1.1; using 10.255.255.254 instead (on interface lo)
26/05/02 00:43:16 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another
address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLe
vel(newLevel).
26/05/02 00:43:17 WARN NativeCodeLoader: Unable to load native-hadoop library f
or your platform... using builtin-java classes where applicable
Spark Version: 3.5.2
```

```
In [2]: pdf = pd.read_excel("data.xlsx")

print(f"Raw data shape: {pdf.shape}")
```

Raw data shape: (541909, 8)

```
In [3]: schema = StructType([
    StructField("InvoiceNo", StringType(), True),
    StructField("StockCode", StringType(), True),
    StructField("Description", StringType(), True),
    StructField("Quantity", IntegerType(), True),
    StructField("InvoiceDate", StringType(), True),
    StructField("UnitPrice", DoubleType(), True),
    StructField("CustomerID", DoubleType(), True),
    StructField("Country", StringType(), True),
])

pdf["Quantity"] = pd.to_numeric(pdf["Quantity"], errors="coerce")
pdf["UnitPrice"] = pd.to_numeric(pdf["UnitPrice"], errors="coerce")
pdf["CustomerID"] = pd.to_numeric(pdf["CustomerID"], errors="coerce")
pdf["InvoiceDate"] = pdf["InvoiceDate"].astype(str)
```

```
df = spark.createDataFrame(pdf, schema=schema)
print(f"Total rows: {df.count()}")
```

26/05/02 00:44:25 WARN TaskSetManager: Stage 0 contains a task of very large size (5154 KiB). The maximum recommended task size is 1000 KiB.

[Stage 0:>

(0 + 8) /

8]

Total rows: 541909

```
In [4]: # Drop no product description
df_clean = df.dropna(subset=["Description"])

# flag cancelled orders
df_clean = df_clean.withColumn("IsCancelled", F.col("InvoiceNo").startswith("C"))

# split valid and cancelled
df_cancelled = df_clean.filter(F.col("IsCancelled") == True)
df_valid = df_clean.filter(F.col("IsCancelled") == False)

# valid only
df_valid = df_valid.filter((F.col("Quantity") > 0) & (F.col("UnitPrice") > 0))

# Cast CustomerID to Integer
df_valid = df_valid.withColumn("CustomerID", F.col("CustomerID").cast(IntegerType))

# Revenue
df_valid = df_valid.withColumn("Revenue", F.round(F.col("Quantity") * F.col("UnitPrice")))
```

```
In [5]: # Parse dates
df_valid = df_valid.withColumn("InvoiceDate", F.to_timestamp("InvoiceDate"))
df_valid = df_valid.filter(F.col("InvoiceDate").isNotNull())
# Extract temporal features
df_valid = (df_valid
    .withColumn("Year", F.year("InvoiceDate"))
    .withColumn("Month", F.month("InvoiceDate"))
    .withColumn("Day", F.dayofmonth("InvoiceDate"))
    .withColumn("DayOfWeek", F.dayofweek("InvoiceDate")) # 1=Sun, 7=Sat
    .withColumn("WeekOfYear", F.weekofyear("InvoiceDate"))
    .withColumn("Quarter", F.quarter("InvoiceDate"))
    .withColumn("Hour", F.hour("InvoiceDate"))
    .withColumn("YearMonth", F.date_format("InvoiceDate", "yyyy-MM"))
)
```

```
In [6]: # Weekend flag
df_valid = df_valid.withColumn(
    "IsWeekend",
    F.when(F.col("DayOfWeek").isin([1, 7]), 1).otherwise(0)
)

print(f"Valid orders: {df_valid.count()}")
print(f"Cancelled orders: {df_cancelled.count()}")
```

```
26/05/02 00:45:01 WARN TaskSetManager: Stage 3 contains a task of very large size (5154 KiB). The maximum recommended task size is 1000 KiB.
```

```
Valid orders:      530104
```

```
26/05/02 00:45:03 WARN TaskSetManager: Stage 6 contains a task of very large size (5154 KiB). The maximum recommended task size is 1000 KiB.
```

```
[Stage 6:>                                     (0 + 8) / 8]
```

```
Cancelled orders: 9288
```

```
In [ ]: # write to HDFS
HDFS_CLEAVED   = "hdfs://localhost:9000/data/retail/cleaved"
HDFS_CANCELLED = "hdfs://localhost:9000/data/retail/cancelled"

df_valid.write.mode("overwrite").partitionBy("YearMonth").parquet(HDFS_CLEAVED)
df_cancelled.write.mode("overwrite").parquet(HDFS_CANCELLED)

print("saved to HDFS")

import subprocess
subprocess.run(["hdfs", "dfs", "-setrep", "3", HDFS_CLEAVED], capture_output=True)
subprocess.run(["hdfs", "dfs", "-setrep", "3", HDFS_CANCELLED], capture_output=True)
print("Replication factor set to 3")

# Verify
print("\n--- HDFS Listing ---")
subprocess.run(["hdfs", "dfs", "-ls", "-R", HDFS_CLEAVED])

print("Done")
spark.stop()
```

Data Analyst

```
In [1]: import findspark
findspark.init()

from pyspark.sql import SparkSession
from pyspark.sql import functions as F
from pyspark.sql.types import *

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use("seaborn-v0_8-whitegrid")
sns.set_palette("husl")

spark = SparkSession.builder \
```

```

.appName("DE_DemandForecasting") \
.master("local[*]") \
.config("spark.sql.adaptive.enabled", "true") \
.config("spark.sql.adaptive.coalescePartitions.enabled", "true") \
.getOrCreate()
print("spark version : ", spark.version)

```

/home/hadoop/.local/lib/python3.10/site-packages/matplotlib/projections/__init__.py:63: UserWarning: Unable to import Axes3D. This may be due to multiple versions of Matplotlib being installed (e.g. as a system package and as a pip package). As a result, the 3D projection is not available.

warnings.warn("Unable to import Axes3D. This may be due to multiple versions of "

26/05/02 01:49:25 WARN Utils: Your hostname, DESKTOP-8HUS7B0 resolves to a loop back address: 127.0.1.1; using 10.255.255.254 instead (on interface lo)

26/05/02 01:49:25 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address

Setting default log level to "WARN".

To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).

26/05/02 01:49:26 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

spark version : 3.5.2

```

In [2]: HDFS_CLEANED    = "hdfs://localhost:9000/data/retail/cleaned"
        HDFS_CANCELLED = "hdfs://localhost:9000/data/retail/cancelled"
        HDFS_OUTPUT    = "hdfs://localhost:9000/data/retail/outputs"

```

```

df = spark.read.parquet(HDFS_CLEANED)
df_cancelled = spark.read.parquet(HDFS_CANCELLED)

print("ok")

```

ok

```

In [4]: df.select("Quantity", "UnitPrice", "Revenue").describe().show()

total_revenue    = df.agg(F.sum("Revenue")).collect()[0][0]
total_orders     = df.select("InvoiceNo").distinct().count()
total_customers  = df.select("CustomerID").distinct().count()
total_products   = df.select("StockCode").distinct().count()
avg_order_value  = df.groupBy("InvoiceNo").agg(F.sum("Revenue").alias("OrderValue"))
avg_order_value  = avg_order_value.agg(F.avg("OrderValue")).collect()[0][0]

print(f"Total Revenue      : {total_revenue:,.2f}")
print(f"Unique Orders      : {total_orders:,}")
print(f"Unique Customers     : {total_customers:,}")
print(f"Unique Products      : {total_products:,}")
print(f"Average Order Value: {avg_order_value:,.2f}")

```

summary	Quantity	UnitPrice	Revenue
count	530104	530104	530104
mean	10.542037034242338	3.907625247122595	20.121871444091383
stddev	155.52412351063685	35.915681104255526	270.3567432189682
min	1	0.001	0.0
max	80995	13541.33	168469.6

Total Revenue : £10,666,684.54
 Unique Orders : 19,960
 Unique Customers : 4,339
 Unique Products : 3,922
 Average Order Value: 534.40

```
In [9]: import matplotlib.dates as mdates

monthly = (df.groupby("Year", "Month")
            .agg(F.round(F.sum("Revenue"), 2).alias("TotalRevenue"))
            .orderBy("Year", "Month")).toPandas()

monthly["YearMonth"] = pd.to_datetime(
    monthly["Year"].astype(str) + "-" + monthly["Month"].astype(str) + "-01"
)

fig, ax = plt.subplots(figsize=(16, 6))

ax.plot(monthly["YearMonth"], monthly["TotalRevenue"],
        marker="o", color="#2E86AB", linewidth=2.5, markersize=8,
        markerfacecolor="white", markeredgewidth=2, markeredgcolor="#2E86AB")


# Format x-axis
ax.xaxis.set_major_formatter(mdates.DateFormatter("%b %Y"))
ax.xaxis.set_major_locator(mdates.MonthLocator(interval=1))
plt.setp(ax.xaxis.get_majorticklabels(), rotation=45, ha="right")

for _, row in monthly.iterrows():
    val = row["TotalRevenue"]
    label = f"£{val/1e6:.2f}M"
    ax.text(row["YearMonth"], val + 40000, label,
            ha="center", va="bottom", fontsize=8, fontweight="bold", color="#3

ax.grid(True, alpha=0.3, linestyle="--")
ax.set_title("Monthly Revenue Over Time", fontsize=15, fontweight="bold", pad=
ax.set_xlabel("Month", fontsize=12)
ax.set_ylabel("Revenue (£)", fontsize=12)
ax.yaxis.set_major_formatter(plt.FuncFormatter(lambda x, _: f"£{x/1e6:.1f}M"))


# Cleaner look
ax.spines["top"].set_visible(False)
ax.spines["right"].set_visible(False)
```

```
plt.tight_layout()
plt.savefig("/tmp/plot_01_monthly_revenue.png", dpi=150, bbox_inches="tight")
plt.show()
```




```
In [10]: # top selling products
top_products = (df.groupby("Description")
                .agg(F.sum("Quantity").alias("TotalQty"))
                .orderBy(F.desc("TotalQty"))
                .limit(10))
tp_pd = top_products.toPandas()

plt.figure(figsize=(12, 6))
plt.barh(tp_pd["Description"][::-1], tp_pd["TotalQty"][::-1], color="coral")
plt.title("Top 10 Best-Selling Products", fontsize=14, fontweight="bold")
plt.xlabel("Total Quantity Sold")
plt.tight_layout()
plt.savefig("/tmp/plot_02_top_products.png", dpi=150)
plt.show()
```



```
In [11]: # by day of week
day_map = {1:"Sun", 2:"Mon", 3:"Tue", 4:"Wed", 5:"Thu", 6:"Fri", 7:"Sat"}
dow = (df.groupby("DayOfWeek")
        .agg(F.sum("Revenue").alias("TotalRevenue"))
        .orderBy("DayOfWeek")).toPandas()
dow["DayName"] = dow["DayOfWeek"].map(day_map)

plt.figure(figsize=(8, 4))
plt.bar(dow["DayName"], dow["TotalRevenue"], color="mediumseagreen")
plt.title("Revenue by Day of Week", fontsize=14, fontweight="bold")
plt.ylabel("Revenue (£)")
plt.tight_layout()
plt.savefig("/tmp/plot_03_dow_revenue.png", dpi=150)
plt.show()
```




```
In [13]: # by country
```

```

countries = (df.groupBy("Country")
              .agg(F.round(F.sum("Revenue"), 2).alias("TotalRevenue"))
              .orderBy(F.desc("TotalRevenue"))
              .limit(10)).toPandas()

plt.figure(figsize=(10, 5))
plt.bar(countries["Country"], countries["TotalRevenue"], color="mediumpurple")
plt.title("Top 10 Countries by Revenue", fontsize=14, fontweight="bold")
plt.ylabel("Revenue (£)")
plt.xticks(rotation=45)
plt.tight_layout()
plt.savefig("/tmp/plot_04_countries.png", dpi=150)
plt.show()

```



```

In [14]: print("Cancelled Order Analysis")

# Cancellation rate overall
cancelled_count = df_cancelled.count()
valid_count     = df.count()
cancel_rate     = cancelled_count / (valid_count + cancelled_count) * 100
print(f"Cancellation Rate: {cancel_rate:.2f}%")

# Revenue lost to cancellations
cancelled_revenue = df_cancelled.withColumn(
    "LostRevenue", F.abs(F.col("Quantity") * F.col("UnitPrice")))
).agg(F.sum("LostRevenue")).collect()[0][0]
print(f"Estimated Revenue Lost to Cancellations: {cancelled_revenue:,.2f}")

# Cancellation rate by country
cancel_by_country = (df_cancelled.groupBy("Country")
                     .agg(F.count("*").alias("CancelledCount")))
valid_by_country  = (df.groupBy("Country")
                     .agg(F.count("*").alias("ValidCount")))

```



```

country_rates = (valid_by_country
    .join(cancel_by_country, "Country", "left_outer")
    .fillna(0)
    .withColumn("CancelRate",
        F.round(F.col("CancelledCount") / (F.col("ValidCount") + F.col("Cancel
    .orderBy(F.desc("CancelRate")))
    .limit(10)).toPandas())

plt.figure(figsize=(10, 5))
plt.bar(country_rates["Country"], country_rates["CancelRate"], color="firebrick")
plt.title("Cancellation Rate by Country (Top 10)", fontsize=14, fontweight="bold")
plt.ylabel("Cancellation Rate (%)")
plt.xticks(rotation=45)
plt.tight_layout()
plt.savefig("/tmp/plot_05_cancel_by_country.png", dpi=150)
plt.show()

# Cancellation by product (top 10 most cancelled)
cancel_by_product = (df_cancelled.groupBy("Description")
    .agg(F.sum(F.abs("Quantity")).alias("CancelledQty"))
    .orderBy(F.desc("CancelledQty"))
    .limit(10)).toPandas()


plt.figure(figsize=(12, 6))
plt.barh(cancel_by_product["Description"][:-1], cancel_by_product["CancelledQty"])
plt.title("Top 10 Most Cancelled Products", fontsize=14, fontweight="bold")
plt.xlabel("Cancelled Quantity")
plt.tight_layout()
plt.savefig("/tmp/plot_06_cancel_by_product.png", dpi=150)
plt.show()


```

Cancelled Order Analysis

Cancellation Rate: 1.72%

Estimated Revenue Lost to Cancellations: £896,812.49





```
In [19]: from datetime import timedelta
from pyspark.sql.window import Window
from pyspark.sql.types import StringType

print("RFM Analysis")

max_date = df.agg(F.max("InvoiceDate")).collect()[0][0]
ref_date = max_date + timedelta(days=1)

rfm = (df.groupBy("CustomerID")
      .agg(
          F.datediff(F.lit(ref_date), F.max("InvoiceDate")).alias("Recency"),
          F.countDistinct("InvoiceNo").alias("Frequency"),
          F.round(F.sum("Revenue"), 2).alias("Monetary")
      ))

# Compute global quantile boundaries
quantiles_r = rfm.approxQuantile("Recency", [0.2, 0.4, 0.6, 0.8], 0.01)
quantiles_f = rfm.approxQuantile("Frequency", [0.2, 0.4, 0.6, 0.8], 0.01)
quantiles_m = rfm.approxQuantile("Monetary", [0.2, 0.4, 0.6, 0.8], 0.01)

# Recency: lower = more recent = better → score 5 for lowest values
@F.udf("int")
def r_score(val):
    if val is None:
        return None
    for i, b in enumerate(sorted(quantiles_r)):
        if val <= b:
            return 5 - i # invert: lowest recency days → score 5
    return 1

# Frequency & Monetary: higher = better → score 5 for highest values
@F.udf("int")
def f_score(val):
```

```

        if val is None:
            return None
        for i, b in enumerate(sorted(quantiles_f)):
            if val <= b:
                return i + 1
        return 5

@F.udf("int")
def m_score(val):
    if val is None:
        return None
    for i, b in enumerate(sorted(quantiles_m)):
        if val <= b:
            return i + 1
    return 5

rfm = (rfm
        .withColumn("R_Score", r_score(F.col("Recency")))
        .withColumn("F_Score", f_score(F.col("Frequency")))
        .withColumn("M_Score", m_score(F.col("Monetary")))
    )

# Segment scoring (max = 15, min = 3)
@F.udf(StringType())
def segment(r, f, m):
    score = r + f + m
    if score >= 13:
        return "TOP"
    elif score >= 10:
        return "Loyal"
    elif score >= 7:
        return "Potential"
    elif score >= 5:
        return "At Risk"
    else:
        return "Lost"

rfm = rfm.withColumn("Segment", segment(F.col("R_Score"), F.col("F_Score"), F.col("M_Score")))

rfm.select("CustomerID", "Recency", "Frequency", "Monetary",
           "R_Score", "F_Score", "M_Score", "Segment").show(10, truncate=False)

# Distribution
segment_dist = (rfm.groupBy("Segment")
                 .agg(
                     F.count("*").alias("Count"),
                     F.round(F.avg("Recency"), 1).alias("AvgRecency"),
                     F.round(F.avg("Frequency"), 1).alias("AvgFrequency"),
                     F.round(F.avg("Monetary"), 2).alias("AvgMonetary")
                 )
                 .orderBy(F.desc("Count"))).toPandas()

print("Segment Distribution:")

```

```

print(segment_dist.to_string(index=False))

# Plot
fig, ax = plt.subplots(figsize=(10, 6))
colors = {"TOP": "#2E86AB", "Loyal": "#A23B72", "Potential": "#F18F01",
          "At Risk": "#C73E1D", "Lost": "#6B6B6B"}
bar_colors = [colors.get(s, "gray") for s in segment_dist["Segment"]]
bars = ax.bar(segment_dist["Segment"], segment_dist["Count"], color=bar_colors)

for bar in bars:
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2., height + 50,
            f'{int(height):,}', ha='center', va='bottom', fontsize=11, fontwei

ax.set_title("Customer Segmentation (RFM)", fontsize=15, fontweight="bold", pa
ax.set_ylabel("Number of Customers", fontsize=12)
ax.set_xlabel("Segment", fontsize=12)
ax.spines["top"].set_visible(False)
ax.spines["right"].set_visible(False)
ax.grid(axis="y", alpha=0.3, linestyle="--")
plt.tight_layout()
plt.savefig("/tmp/plot_rfm_segments.png", dpi=150, bbox_inches="tight")
plt.show()

rfm.write.mode("overwrite").parquet(f"{HDFS_OUTPUT}/rfm_segments")
print("RFM saved to HDFS")

```

RFM Analysis

```


+-----+-----+-----+-----+-----+-----+-----+
|CustomerID|Recency|Frequency|Monetary|R_Score|F_Score|M_Score|Segment|
+-----+-----+-----+-----+-----+-----+-----+
|15619      |11      |1         |336.4   |5        |1        |2        |Potential|
|15727      |17      |7         |5178.96 |4        |5        |5        |TOP       |
|15790      |11      |1         |220.85  |5        |1        |1        |Potential|
|13832      |20      |1         |52.2    |4        |1        |1        |At Risk   |
|13623      |31      |5         |747.78  |4        |4        |3        |Loyal     |
|15957      |32      |1         |428.89  |4        |1        |2        |Potential|
|16503      |107     |4         |1431.93 |2        |4        |4        |Loyal     |
|17389      |1       |34        |31833.68|5        |5        |5        |TOP       |
|17420      |51      |3         |598.83  |3        |3        |3        |Potential|
|16386      |29      |2         |317.2   |4        |2        |2        |Potential|
+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 10 rows

Segment Distribution:

Segment	Count	AvgRecency	AvgFrequency	AvgMonetary
Potential	1058	75.0	2.2	893.36
TOP	931	14.9	13.3	8397.50
Loyal	842	44.1	4.1	1728.79
At Risk	789	122.9	1.3	381.17
Lost	719	245.3	1.0	204.53



RFM saved to HDFS

```
In [20]: print("Retention Analysis")

# first purchase month per customer (cohort)
customer_cohort = (df.groupBy("CustomerID")
                  .agg(F.min("YearMonth").alias("CohortMonth")))

df_cohort = df.join(customer_cohort, "CustomerID")

# Period number = months since first purchase
df_cohort = df_cohort.withColumn(
    "PeriodNumber",
    (F.year("YearMonth") * 12 + F.month("YearMonth")) -
    (F.year("CohortMonth") * 12 + F.month("CohortMonth"))
)

# Unique customers per cohort & period
cohort_data = (df_cohort.groupBy("CohortMonth", "PeriodNumber")
              .agg(F.countDistinct("CustomerID").alias("CustomerCount")))

# Cohort sizes
cohort_sizes = (customer_cohort.groupBy("CohortMonth")
               .agg(F.count("*").alias("CohortSize")))

cohort_table = cohort_data.join(cohort_sizes, "CohortMonth")
cohort_table = cohort_table.withColumn(
    "RetentionRate",
    F.round(F.col("CustomerCount") / F.col("CohortSize") * 100, 2)
)

# Pivot to matrix
```

```


cohort_matrix = (cohort_table
    .groupBy("CohortMonth")
    .pivot("PeriodNumber")
    .agg(F.first("RetentionRate"))
    .orderBy("CohortMonth")).toPandas()

# Plot heatmap
cohort_matrix_indexed = cohort_matrix.set_index("CohortMonth")
plt.figure(figsize=(14, 8))
sns.heatmap(cohort_matrix_indexed.iloc[:, :13], # first 13 periods
    annot=True, fmt=".1f", cmap="YlGnBu", linewidths=0.5)
plt.title("Cohort Retention Matrix (%)", fontsize=14, fontweight="bold")
plt.ylabel("Cohort Month")
plt.xlabel("Period Number (Months)")
plt.tight_layout()
plt.savefig("/tmp/plot_08_cohort_retention.png", dpi=150)
plt.show()

cohort_table.write.mode("overwrite").parquet(f"{HDFS_OUTPUT}/cohort_retention")
print("saved to HDFS")

```

Retention Analysis



saved to HDFS

```

In [22]: print("Retention Analysis")

# First purchase + total orders per customer
customer_stats = (df.groupBy("CustomerID")
    .agg(

```

```

        F.min("InvoiceDate").alias("FirstPurchase"),
        F.max("InvoiceDate").alias("LastPurchase"),
        F.countDistinct("InvoiceNo").alias("TotalOrders")
    ))

customer_stats = customer_stats.withColumn(
    "IsRetained", (F.col("TotalOrders") > 1).cast("int")
)

# Overall numbers
total      = customer_stats.count()
retained   = customer_stats.filter(F.col("IsRetained") == 1).count()
one_time   = total - retained
rate       = round(retained / total * 100, 2)

print(f"Total Customers : {total:,}")
print(f"Retained          : {retained:,}   ({rate}%)" )
print(f"One-time only    : {one_time:,}   ({round(100 - rate, 2)}%)" )

# Return frequency distribution (how many customers made exactly N purchases)
order_dist = (customer_stats
    .groupBy("TotalOrders")
    .agg(F.count("*").alias("CustomerCount"))
    .orderBy("TotalOrders")
    .toPandas())

order_dist["RetentionRate"] = (
    order_dist["CustomerCount"][::-1].cumsum()[::-1] / total * 100
).round(2)

# Plot
fig, ax1 = plt.subplots(figsize=(13, 6))

# Cap display at 15 orders (long tail is noise)
plot_df = order_dist[order_dist["TotalOrders"] <= 15].copy()

# Bar – number of customers per order count
ax1.bar(plot_df["TotalOrders"], plot_df["CustomerCount"],
        color="#2E86AB", edgecolor="white", linewidth=1.2, label="# Customers"

for _, row in plot_df.iterrows():
    ax1.text(row["TotalOrders"], row["CustomerCount"] + 10,
            f'{int(row["CustomerCount"]):,}',
            ha="center", va="bottom", fontsize=8.5, fontweight="bold", color=

ax1.set_xlabel("Number of Purchases", fontsize=12)
ax1.set_ylabel("Number of Customers", fontsize=12, color="#2E86AB")
ax1.tick_params(axis="y", labelcolor="#2E86AB")
ax1.spines["top"].set_visible(False)
ax1.set_xticks(plot_df["TotalOrders"])

# Line – cumulative retention rate (% of customers who made at least N purchas
ax2 = ax1.twinx()
```

```

ax2.plot(plot_df["TotalOrders"], plot_df["RetentionRate"],
        color="#C73E1D", linewidth=2.5, marker="o", markersize=6, label="Retention Rate")

for _, row in plot_df.iterrows():
    ax2.text(row["TotalOrders"], row["RetentionRate"] + 1.5,
            f'{row["RetentionRate"]:.1f}%',
            ha="center", fontsize=8, color="#C73E1D", fontweight="bold")

ax2.set_ylabel("Cumulative Retention Rate (%)", fontsize=12, color="#C73E1D")
ax2.tick_params(axis="y", labelcolor="#C73E1D")
ax2.spines["top"].set_visible(False)
ax2.set_ylim(0, 110)

# Legends
lines1, labels1 = ax1.get_legend_handles_labels()
lines2, labels2 = ax2.get_legend_handles_labels()
ax1.legend(lines1 + lines2, labels1 + labels2, loc="upper right", fontsize=10)

plt.title("Customer Purchase Frequency & Cumulative Retention Rate",
        fontsize=14, fontweight="bold", pad=15)
plt.tight_layout()
plt.savefig("/tmp/plot_retention.png", dpi=150, bbox_inches="tight")
plt.show()


```

Retention Analysis

Total Customers : 4,339

Retained : 2,846 (65.59%)

One-time only : 1,493 (34.41%)



```
In [8]: print("Product Portfolio Analysis")
```

```

# Revenue + quantity per product
product_stats = (df.groupby("StockCode", "Description")
    .agg(
        F.round(F.sum("Revenue"), 2).alias("ProductRevenue"),
        F.sum("Quantity").alias("TotalQuantity"),
        F.countDistinct("InvoiceNo").alias("TotalOrders")
    )

```



```

))

# Compute medians as split threshold
revenue_median = product_stats.approxQuantile("ProductRevenue", [0.5], 0.01)
quantity_median = product_stats.approxQuantile("TotalQuantity", [0.5], 0.01)

print(f"Revenue median : {revenue_median:,.2f}")
print(f"Quantity median : {quantity_median:,.0f} units")

# Classify into quadrants
product_stats = product_stats.withColumn(
    "Quadrant",
    F.when((F.col("ProductRevenue") >= revenue_median) & (F.col("TotalQuantity"
        .when((F.col("ProductRevenue") >= revenue_median) & (F.col("TotalQuantity"
        .when((F.col("ProductRevenue") < revenue_median) & (F.col("TotalQuantity"
        .otherwise("Low Revenue / Low Volume")
    )
)

# Summary per quadrant
quadrant_summary = (product_stats.groupBy("Quadrant")
    .agg(
        F.count("*").alias("ProductCount"),
        F.round(F.sum("ProductRevenue"), 2).alias("TotalRevenue"),
        F.round(F.avg("ProductRevenue"), 2).alias("AvgRevenue"),
        F.round(F.avg("TotalQuantity"), 0).alias("AvgQuantity")
    )
    .orderBy(F.desc("TotalRevenue")))

print("\nQuadrant Summary:")
quadrant_summary.show(truncate=False)

# Top 5 per quadrant
print("\nTop 5 per Quadrant (by Revenue):")
for quadrant in ["High Revenue / High Volume", "High Revenue / Low Volume",
    "Low Revenue / High Volume", "Low Revenue / Low Volume"]:
    print(f"\n-- {quadrant} --")
    (product_stats
        .filter(F.col("Quadrant") == quadrant)
        .orderBy(F.desc("ProductRevenue"))
        .select("StockCode", "Description", "ProductRevenue", "TotalQuantity",
        .limit(5)
        .show(truncate=False))

# --- Plot ---
plot_df = product_stats.sample(fraction=0.8).limit(500).toPandas()
quad_pd = quadrant_summary.toPandas()

colors = {
    "High Revenue / High Volume": "#2E86AB",
    "High Revenue / Low Volume": "#A23B72",
    "Low Revenue / High Volume": "#F18F01",
    "Low Revenue / Low Volume": "#AAAAAA"
}

```

```

fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Scatter – revenue vs quantity
for quadrant, group in plot_df.groupby("Quadrant"):
    axes[0].scatter(group["TotalQuantity"], group["ProductRevenue"],
                    label=quadrant, color=colors[quadrant],
                    alpha=0.7, s=50, edgecolors="white")
axes[0].axvline(quantity_median, color="gray", linestyle="--", linewidth=1)
axes[0].axhline(revenue_median, color="gray", linestyle="--", linewidth=1)
axes[0].set_xlabel("Total Quantity Sold")
axes[0].set_ylabel("Total Revenue (£)")
axes[0].set_title("Product Portfolio Matrix", fontweight="bold", fontsize=13)
axes[0].legend(fontsize=8)
axes[0].spines["top"].set_visible(False)
axes[0].spines["right"].set_visible(False)

# Bar – product count per quadrant
bar_colors = [colors[q] for q in quad_pd["Quadrant"]]
bars = axes[1].bar(quad_pd["Quadrant"], quad_pd["ProductCount"],
                  color=bar_colors, edgecolor="white", linewidth=1.5)
for bar, rev in zip(bars, quad_pd["TotalRevenue"]):
    axes[1].text(bar.get_x() + bar.get_width()/2, bar.get_height() + 5,
                f'£{rev:,.0f}', ha="center", fontsize=9, fontweight="bold")
axes[1].set_title("Products per Quadrant\n(label = total revenue)", fontweight="bold")
axes[1].set_ylabel("Number of Products")
axes[1].tick_params(axis="x", labelsize=8)
axes[1].spines["top"].set_visible(False)
axes[1].spines["right"].set_visible(False)

plt.tight_layout()
plt.savefig("/tmp/plot_product_portfolio.png", dpi=150, bbox_inches="tight")
plt.show()

product_stats.write.mode("overwrite").parquet(f"{HDFS_OUTPUT}/product_portfolio")
print("Saved to HDFS")

```

Product Portfolio Analysis
Revenue median : £607.73
Quantity median : 315 units

Quadrant Summary:

Quadrant	ProductCount	TotalRevenue	AvgRevenue	AvgQuantity
High Revenue / High Volume	1803	9829067.84	5451.51	2871.0
High Revenue / Low Volume	300	472911.15	1576.37	195.0
Low Revenue / Low Volume	1758	247995.44	141.07	64.0
Low Revenue / High Volume	300	116710.11	389.03	802.0

Top 5 per Quadrant (by Revenue):

-- High Revenue / High Volume --

StockCode	Description	ProductRevenue	TotalQuantity	TotalOrders
DOT	DOTCOM POSTAGE	206248.77	706	706
22423	REGENCY CAKESTAND 3 TIER	174484.74	13879	1988
23843	PAPER CRAFT , LITTLE BIRDIE	168469.6	80995	1
85123A	WHITE HANGING HEART T-LIGHT HOLDER	104340.29	37599	2189
47566	PARTY BUNTING	99504.33	18295	1685

-- High Revenue / Low Volume --

+-----+-----+-----+-----+				
+-----+				
StockCode	Description	ProductRevenue	TotalQuantity	TotalOrders
+-----+-----+-----+-----+				
+-----+				
22502	PICNIC BASKET WICKER 60 PIECES	39619.5	61	2
AMAZONFEE	AMAZON FEE	13761.09	2	2
B	Adjust bad debt	11062.06	1	1
22655	VINTAGE RED KITCHEN CABINET	8125.0	60	38
C2	CARRIAGE	7051.0	142	141
+-----+-----+-----+-----+				
+-----+				



-- Low Revenue / High Volume --

+-----+-----+-----+-----+				
+-----+				
StockCode	Description	ProductRevenue	TotalQuantity	TotalOrders
+-----+-----+-----+-----+				
+-----+				
72598	SET/12 TAPER CANDLES	605.89	607	93
22595	CHRISTMAS GINGHAM HEART	605.8	649	54
22999	TRAVEL CARD WALLET VINTAGE LEAF	605.78	1415	117
84625A	PINK NEW BAROQUECANDLESTICK CANDLE	605.48	478	42
22232	JIGSAW TOADSTOOLS 3 PIECE	604.46	690	83
+-----+-----+-----+-----+				
+-----+				

-- Low Revenue / Low Volume --

+-----+-----+-----+-----+				
+-----+				
StockCode	Description	ProductRevenue	TotalQuantity	TotalOrders
+-----+-----+-----+-----+				
+-----+				
21769	VINTAGE POST OFFICE CABINET	607.65	11	2
21709	FOLDING UMBRELLA CHOCOLATE POLKADOT	607.23	133	59
23520	EMBROIDERED RIBBON REEL RUBY	604.3	168	32

22642	SET OF 4 NAPKIN CHARMS STARS	596.36	269	51
21253	SET OF PICTURE FRAME STICKERS	593.37	135	62
+-----+-----+-----+-----+				
+-----+				



[Stage 222:=====>
7]

(3 + 4) /

Saved to HDFS

ML Engineer

```
In [19]: import findspark
findspark.init()

from pyspark.sql import SparkSession
from pyspark.sql import functions as F
from pyspark.sql.types import *
from pyspark.sql.window import Window
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.regression import LinearRegression, RandomForestRegressor, GBT
from pyspark.ml.evaluation import RegressionEvaluator
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

spark = SparkSession.builder \
    .appName("ML_DemandForecasting") \
    .master("local[*]") \
    .getOrCreate()

HDFS_CLEANED = "hdfs://localhost:9000/data/retail/cleaned"
HDFS_OUTPUT = "hdfs://localhost:9000/data/retail/outputs"
HDFS_MODELS = "hdfs://localhost:9000/data/retail/models"
```

```
In [20]: # aggregate weekly demand
df = spark.read.parquet(HDFS_CLEANED)

weekly = (df.groupBy("Year", "WeekOfYear", "StockCode")
          .agg(
            F.sum("Quantity").alias("Demand"),
            F.sum("Revenue").alias("Revenue"),
            F.countDistinct("InvoiceNo").alias("NumOrders"),
            F.countDistinct("CustomerID").alias("NumCustomers"),
            F.avg("UnitPrice").alias("AvgPrice")
          )
          .orderBy("StockCode", "Year", "WeekOfYear"))

print(f"rows: {weekly.count()}")
weekly.show(5)
```

rows: 101579

[Stage 722:> (0 + 8) / 8]

```
+---+-----+-----+-----+-----+-----+-----+
+-----+
|Year|WeekOfYear|StockCode|Demand|Revenue|NumOrders|NumCustomers|
AvgPrice|
+---+-----+-----+-----+-----+-----+-----+
+-----+
|2010|      48|   10002|    70|65.16999999999999|      7|      5|1.08
14285714285714|
|2010|      49|   10002|   138|    127.02|     14|      9|  1.1
97142857142857|
|2010|      50|   10002|    41|38.89999999999999|      8|      5|
1.255|
|2010|      51|   10002|     2|     3.32|      1|      1|
1.66|
|2011|       1|   10002|    73|    62.86|      3|      3|1.11
99999999999999|
+---+-----+-----+-----+-----+-----+-----+
+-----+
only showing top 5 rows
```

```
In [21]: print("Feature Engineering")

# Fix types at the root
weekly = (weekly
          .withColumn("Demand", F.col("Demand").cast("double"))
          .withColumn("AvgPrice", F.col("AvgPrice").cast("double"))
          .withColumn("WeekOfYear", F.col("WeekOfYear").cast("double"))
          .withColumn("Year", F.col("Year").cast("double"))
          .withColumn("NumOrders", F.col("NumOrders").cast("double"))
          .withColumn("NumCustomers", F.col("NumCustomers").cast("double"))
          )
```

```

product_window = Window.partitionBy("StockCode").orderBy("Year", "WeekOfYear")

# Lag features
weekly = (weekly
    .withColumn("Lag1", F.lag("Demand", 1).over(product_window))
    .withColumn("Lag2", F.lag("Demand", 2).over(product_window))
    .withColumn("Lag4", F.lag("Demand", 4).over(product_window))
    .withColumn("Lag8", F.lag("Demand", 8).over(product_window))
)

# Rolling statistics
weekly = (weekly
    .withColumn("RollingMean4", F.avg("Demand").over(product_window.rowsBetween(-4, 4)))
    .withColumn("RollingMean8", F.avg("Demand").over(product_window.rowsBetween(-8, 8)))
    .withColumn("RollingStd4", F.stddev("Demand").over(product_window.rowsBetween(-4, 4)))
)

# Seasonality
weekly = (weekly
    .withColumn("WeekSin", F.sin(2 * np.pi * F.col("WeekOfYear") / 52))
    .withColumn("WeekCos", F.cos(2 * np.pi * F.col("WeekOfYear") / 52))
)

# Price features
weekly = weekly.withColumn("PriceLag1", F.lag("AvgPrice", 1).over(product_window))

# Drop nulls from lags
weekly = weekly.dropna()
weekly = weekly.withColumn("Demand", F.col("Demand").cast("double")) # ensure double

print(f"Rows after feature eng: {weekly.count()}")
weekly.printSchema() # verify all doubles

```

```

[Stage 727:=====> (183 + 9) / 200]
0]

```

Feature Engineering

```

[Stage 769:=====> (181 + 8) / 200][Stage 776:> (0 + 0) / 1]
1]

```

Rows after feature eng: 73723

```
root
|-- Year: double (nullable = true)
|-- WeekOfYear: double (nullable = true)
|-- StockCode: string (nullable = true)
|-- Demand: double (nullable = true)
|-- Revenue: double (nullable = true)
|-- NumOrders: double (nullable = false)
|-- NumCustomers: double (nullable = false)
|-- AvgPrice: double (nullable = true)
|-- Lag1: double (nullable = true)
|-- Lag2: double (nullable = true)
|-- Lag4: double (nullable = true)
|-- Lag8: double (nullable = true)
|-- RollingMean4: double (nullable = true)
|-- RollingMean8: double (nullable = true)
|-- RollingStd4: double (nullable = true)
|-- WeekSin: double (nullable = true)
|-- WeekCos: double (nullable = true)
|-- PriceLag1: double (nullable = true)
```

```
In [22]: # split time based
# Train: Dec 2010 → Sep 2011
# Test: Oct 2011 → Dec 2011

# clac expanding mean along

train_df = weekly.filter(
    (F.col("Year") == 2010) |
    ((F.col("Year") == 2011) & (F.col("WeekOfYear") <= 39))
).withColumn("ExpandingMean",
    F.avg("Demand").over(product_window.rowsBetween(Window.unboundedPreceding,

# last known ExpandingMean from train per product
last_expanding = (train_df
    .groupBy("StockCode")
    .agg(F.last("ExpandingMean", ignorenulls=True).alias("ExpandingMean")))

test_df = (weekly
    .filter((F.col("Year") == 2011) & (F.col("WeekOfYear") > 39))
    .drop("ExpandingMean")
    .join(last_expanding, "StockCode", "left"))
print(train_df.count(), test_df.count())
```

```
[Stage 853:=====> (154 + 8) / 200][Stage 860:> (0 + 0) /
1]
52931 20792
```

```
In [23]: # historical product features
product_profile = (train_df.groupBy("StockCode")
```



```

    .agg(
        F.avg("Demand").alias("ProductAvgDemand"),
        F.stddev("Demand").alias("ProductDemandStd"),
        ((F.last("Demand") - F.first("Demand")) / F.count("*")).alias("ProductTrend")
    )
    .fillna(0))

train_df = train_df.join(product_profile, "StockCode", "left")
test_df = test_df.join(product_profile, "StockCode", "left")

# # Encode StockCode
# from pyspark.ml.feature import StringIndexer
# indexer_model = StringIndexer(inputCol="StockCode", outputCol="StockCodeIdx")
# train_df = indexer_model.transform(train_df)
# test_df = indexer_model.transform(test_df)

```

```

In [24]: feature_cols = [
#     "StockCodeIdx",
    "ProductAvgDemand", "ProductDemandStd", "ProductTrend",
    "WeekOfYear", "Year", "WeekSin", "WeekCos",
    "NumOrders", "NumCustomers", "AvgPrice", "PriceLag1",
    "Lag1", "Lag2", "Lag4", "Lag8",
    "RollingMean4", "RollingMean8", "RollingStd4",
    "ExpandingMean"
]

assembler = VectorAssembler(inputCols=feature_cols, outputCol="features", handleInvalid="skip")
train_vec = assembler.transform(train_df).select("features", F.col("Demand").alias("label"))
test_vec = assembler.transform(test_df).select("features", F.col("Demand").alias("label"))

train_vec.cache()
test_vec.cache()

```

```

[Stage 865:=====>(197 + 3) / 200]

```

```

Out[24]: DataFrame[features: vector, label: double, StockCode: string, Year: double, WeekOfYear: double]

```

```

In [25]: def evaluate(df_pred, name):
    evaluator = RegressionEvaluator(labelCol="label", predictionCol="prediction")
    rmse = evaluator.setMetricName("rmse").evaluate(df_pred)
    mae = evaluator.setMetricName("mae").evaluate(df_pred)
    r2 = evaluator.setMetricName("r2").evaluate(df_pred)
    mape = (df_pred
        .filter(F.col("label") > 0)
        .withColumn("pct", F.abs((F.col("label") - F.col("prediction")) / F.col("label")))
        .agg(F.avg("pct").collect()[0][0]) * 100)
    print(f"{name:<30} RMSE:{round(rmse,2):>8} MAE:{round(mae,2):>8} MAPE:{round(mape,2):>8}")
    return {"RMSE": round(rmse,2), "MAE": round(mae,2), "MAPE": round(mape,2)}

results = {}

# baselines

```

```

print("Baselines")
results["Naive (Lag1)"] = evaluate(
    test_df.withColumn("prediction", F.col("Lag1"))
    .select(F.col("Demand").alias("label"), "prediction"), "Naive (Lag1)

results["MA-4"] = evaluate(
    test_df.withColumn("prediction", F.col("RollingMean4"))
    .select(F.col("Demand").alias("label"), "prediction"), "MA-4")

print("Linear Regression")
from pyspark.ml.regression import LinearRegression
lr = LinearRegression(featuresCol="features", labelCol="label",
    maxIter=100, regParam=0.1, elasticNetParam=0.0)
lr_model = lr.fit(train_vec)
lr_preds = lr_model.transform(test_vec)
results["Linear Regression"] = evaluate(lr_preds, "Linear Regression")

print("Random Forest")
from pyspark.ml.regression import RandomForestRegressor
rf = RandomForestRegressor(featuresCol="features", labelCol="label",
    numTrees=20, maxDepth=4, seed=42, maxBins=3000)
rf_model = rf.fit(train_vec)
rf_preds = rf_model.transform(test_vec)
results["Random Forest"] = evaluate(rf_preds, "Random Forest")

print("Gradient Boosting")
from pyspark.ml.regression import GBRegressor
gbt = GBRegressor(featuresCol="features", labelCol="label",
    maxIter=20, maxDepth=4, stepSize=0.1, seed=42, maxBins=3000)
gbt_model = gbt.fit(train_vec)
gbt_preds = gbt_model.transform(test_vec)
results["Gradient Boosting"] = evaluate(gbt_preds, "Gradient Boosting")

```

Baselines

Naive (Lag1)	RMSE: 144.07	MAE: 50.34	MAPE:204.81%	R2: 0.2951
MA-4	RMSE: 121.47	MAE: 44.76	MAPE:205.98%	R2: 0.4989
Linear Regression				
Linear Regression	RMSE: 131.81	MAE: 49.44	MAPE:313.76%	R2: 0.4409
Random Forest				
Random Forest	RMSE: 137.93	MAE: 43.81	MAPE:229.02%	R2: 0.3878
Gradient Boosting				
Gradient Boosting	RMSE: 126.72	MAE: 41.48	MAPE:156.16%	R2: 0.4833

```

In [26]: print(f"{'Model':<30} {'RMSE':>8} {'MAE':>8} {'MAPE':>8} {'R2':>8}")
for model, m in results.items():
    r2 = str(m["R2"]) if m["R2"] is not None else "N/A"
    print(f"{model:<30} {str(m['RMSE']):>8} {str(m['MAE']):>8} {str(m['MAPE'])

fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# Get predictions from all models
models_to_plot = {
    "Linear Regression": lr_preds,
    "Random Forest": rf_preds,
    "Gradient Boosting": gbt_preds
}

for idx, (model_name, predictions) in enumerate(models_to_plot.items()):
    row, col = divmod(idx, 2)

    # Sample 200 points for visualization
    plot_df = predictions.select("label", "prediction").toPandas().head(200)

    axes[row, col].plot(plot_df["label"].values, label="Actual", color="#2E86A
    axes[row, col].plot(plot_df["prediction"].values, label=f"Predicted", colc
    axes[row, col].set_title(f"{model_name}\n(RMSE: {results[model_name]['RMSE
        fontweight="bold", fontsize=12)
    axes[row, col].set_xlabel("Sample Index")
    axes[row, col].set_ylabel("Weekly Demand")
    axes[row, col].legend()
    axes[row, col].spines["top"].set_visible(False)
    axes[row, col].spines["right"].set_visible(False)

plt.suptitle("Actual vs Predicted Demand - Model Comparison", fontweight="bold
plt.tight_layout()

# residual analysis
fig2, axes2 = plt.subplots(1, 3, figsize=(15, 5))

for idx, (model_name, predictions) in enumerate(models_to_plot.items()):
    # Calculate residuals
    resid_df = predictions.select(
        "label", "prediction",
        (F.col("label") - F.col("prediction")).alias("residual")
    ).toPandas()

    # Plot residuals
    axes2[idx].scatter(resid_df["prediction"], resid_df["residual"], alpha=0.5
    axes2[idx].axhline(y=0, color='red', linestyle='--', linewidth=1)
    axes2[idx].set_title(f"{model_name}", fontweight="bold", fontsize=12)
    axes2[idx].set_xlabel("Predicted Demand")
    axes2[idx].set_ylabel("Residuals (Actual - Predicted)")
    axes2[idx].spines["top"].set_visible(False)
    axes2[idx].spines["right"].set_visible(False)

plt.suptitle("Residual Plot Analysis", fontweight="bold", fontsize=14)

```

```

plt.tight_layout()

# feature importance
fig3, axes3 = plt.subplots(1, 2, figsize=(14, 8))

# Random Forest Feature Importance
rf_importances = rf_model.featureImportances.toArray()
rf_feat_imp = (pd.DataFrame({"Feature": feature_cols, "Importance": rf_importances
                             .sort_values("Importance", ascending=True)
                             .tail(15)}))

axes3[0].barh(rf_feat_imp["Feature"], rf_feat_imp["Importance"], color="#2E86A1")
axes3[0].set_title("Random Forest - Feature Importance (Top 15)", fontweight="bold")
axes3[0].set_xlabel("Importance")
axes3[0].spines["top"].set_visible(False)
axes3[0].spines["right"].set_visible(False)

# Gradient Boosting Feature Importance
gbt_importances = gbt_model.featureImportances.toArray()
gbt_feat_imp = (pd.DataFrame({"Feature": feature_cols, "Importance": gbt_importances
                              .sort_values("Importance", ascending=True)
                              .tail(15)}))

axes3[1].barh(gbt_feat_imp["Feature"], gbt_feat_imp["Importance"], color="#A23B72")
axes3[1].set_title("Gradient Boosting - Feature Importance (Top 15)", fontweight="bold")
axes3[1].set_xlabel("Importance")
axes3[1].spines["top"].set_visible(False)
axes3[1].spines["right"].set_visible(False)

plt.suptitle("Feature Importance Comparison", fontweight="bold", fontsize=14)
plt.tight_layout()

fig4, ax4 = plt.subplots(figsize=(10, 6))

error_data = []
model_names = []
for model_name, predictions in models_to_plot.items():
    errors = predictions.select(
        (F.abs(F.col("label") - F.col("prediction"))).alias("abs_error")
    ).toPandas()["abs_error"].values
    error_data.append(errors)
    model_names.append(model_name)

bp = ax4.boxplot(error_data, labels=model_names, patch_artist=True)
for patch, color in zip(bp['boxes'], ['#2E86A1', '#A23B72', '#C73E1D']):
    patch.set_facecolor(color)
    patch.set_alpha(0.7)

ax4.set_title("Absolute Error Distribution by Model", fontweight="bold", fontstyle="italic")
ax4.set_ylabel("Absolute Error")
ax4.set_yscale('log') # Log scale if errors have wide range
ax4.spines["top"].set_visible(False)
ax4.spines["right"].set_visible(False)

```

```

ax4.grid(axis='y', alpha=0.3)

plt.tight_layout()

fig5, axes5 = plt.subplots(1, 3, figsize=(15, 5))

for idx, (model_name, predictions) in enumerate(models_to_plot.items()):
    plot_df = predictions.select("label", "prediction").toPandas()

    axes5[idx].scatter(plot_df["label"], plot_df["prediction"], alpha=0.3, s=5)

    # Perfect prediction line
    max_val = max(plot_df["label"].max(), plot_df["prediction"].max())
    axes5[idx].plot([0, max_val], [0, max_val], 'r--', alpha=0.5, label="Perfect")

    axes5[idx].set_title(f"{model_name}\nR2 = {results[model_name]['R2']}", fontweight="bold")
    axes5[idx].set_xlabel("Actual Demand")
    axes5[idx].set_ylabel("Predicted Demand")
    axes5[idx].set_xlim(0, max_val)
    axes5[idx].set_ylim(0, max_val)
    axes5[idx].legend()
    axes5[idx].spines["top"].set_visible(False)
    axes5[idx].spines["right"].set_visible(False)

plt.suptitle("Actual vs Predicted - Scatter Plot Analysis", fontweight="bold", y=1.05)
plt.tight_layout()

try:
    lr_model.write().overwrite().save(f"{HDFS_MODELS}/lr_model")
except NameError:
    print("Linear Regression model not available for saving")

summary = pd.DataFrame(results).T
summary.to_csv("/tmp/model_summary.csv")

print("Summary:")
print(summary.round(2).to_string())

```

Model	RMSE	MAE	MAPE	R2
Naive (Lag1)	144.07	50.34	204.81%	0.2951
MA-4	121.47	44.76	205.98%	0.4989
Linear Regression	131.81	49.44	313.76%	0.4409
Random Forest	137.93	43.81	229.02%	0.3878
Gradient Boosting	126.72	41.48	156.16%	0.4833



/tmp/ipykernel_10629/4035172233.py:97: MatplotlibDeprecationWarning: The 'label s' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.

```
bp = ax4.boxplot(error_data, labels=model_names, patch_artist=True)
```


Summary:

	RMSE	MAE	MAPE	R2
Naive (Lag1)	144.07	50.34	204.81	0.30
MA-4	121.47	44.76	205.98	0.50
Linear Regression	131.81	49.44	313.76	0.44
Random Forest	137.93	43.81	229.02	0.39
Gradient Boosting	126.72	41.48	156.16	0.48


Actual vs Predicted Demand - Model Comparison




Feature Importance Comparison



Absolute Error Distribution by Model





In [27]: `spark.stop()`

Big Data Eng

In [28]: `from pyspark.sql import SparkSession`
`from pyspark.sql import functions as F`
`import time`

```
spark = SparkSession.builder \
    .appName("BD_DemandForecasting") \
    .master("local[*]") \
    .getOrCreate()
```

```
HDFS_CLEANED = "hdfs://localhost:9000/data/retail/cleaned"
```

```
df = spark.read.parquet(HDFS_CLEANED)
print(f"Loaded {df.count()} rows")
```

```
[Stage 1:=====>
7]
```

(2 + 5) /

Loaded 530104 rows

In [29]: `print(f"Partitions in raw df: {df.rdd.getNumPartitions()}")`

```
# Repartition for product-level aggregations
```

```
df_repart = df.repartition(8, "StockCode")
```

```
print(f"Partitions after repartition(8, StockCode): {df_repart.rdd.getNumPartitions()}")
```

Partitions in raw df: 7

```
[Stage 4:=====>
7]
```

(6 + 1) /

Partitions after repartition(8, StockCode): 8

In [30]: `# Without repartition`

```
start = time.time()
```

```
result1 = df.groupBy("StockCode").agg(F.sum("Quantity").alias("TotalQty")).collect()
```



```

t1 = time.time() - start
print(f"Without repartition: {t1:.2f}s")

# With repartition
start = time.time()
result2 = df_repart.groupBy("StockCode").agg(F.sum("Quantity").alias("TotalQty"))
t2 = time.time() - start
print(f"With repartition: {t2:.2f}s")
print(f"Speedup: {t1/t2:.2f}x")

```

Without repartition: 2.29s

[Stage 8:=====>
7]

(3 + 4) /

With repartition: 1.85s

Speedup: 1.23x

```

In [31]: # Before cache
start = time.time()
df.filter(F.col("Country") == "United Kingdom").count()
df.filter(F.col("Country") == "United Kingdom").count()
t_no_cache = time.time() - start
print(f"Without cache : {t_no_cache:.2f}s")

# After cache
df_cached = df.cache()
df_cached.count()

start = time.time()
df_cached.filter(F.col("Country") == "United Kingdom").count()
df_cached.filter(F.col("Country") == "United Kingdom").count()
t_cache = time.time() - start
print(f"With cache: {t_cache:.2f}s")
print(f"Cache speedup: {t_no_cache/t_cache:.2f}x")

```

Without cache : 3.20s

With cache: 0.82s

Cache speedup: 3.89x

```

In [32]: print(f"Partitions before coalesce: {df.rdd.getNumPartitions()}")
df_coalesced = df.coalesce(4)
print(f"Partitions after coalesce(4): {df_coalesced.rdd.getNumPartitions()}")

# Write optimized output
HDFS_OPTIMIZED = "hdfs://localhost:9000/data/retail/optimized"
df_coalesced.write.mode("overwrite").parquet(HDFS_OPTIMIZED)
print(f"Coalesced output written to {HDFS_OPTIMIZED}")

```

Partitions before coalesce: 7

Partitions after coalesce(4): 4

Coalesced output written to hdfs://localhost:9000/data/retail/optimized

```
In [33]: print("Execution Plan (groupBy StockCode)")
df.groupBy("StockCode").agg(F.sum("Quantity").alias("TotalQty")).explain()
```

```
Execution Plan (groupBy StockCode)
== Physical Plan ==
AdaptiveSparkPlan isFinalPlan=false
+- HashAggregate(keys=[StockCode#38492], functions=[sum(Quantity#38494)])
   +- Exchange hashpartitioning(StockCode#38492, 200), ENSURE_REQUIREMENTS, [plan_id=33478]
      +- HashAggregate(keys=[StockCode#38492], functions=[partial_sum(Quantity#38494)])
         +- InMemoryTableScan [StockCode#38492, Quantity#38494]
            +- InMemoryRelation [InvoiceNo#38491, StockCode#38492, Description#38493, Quantity#38494, InvoiceDate#38495, UnitPrice#38496, CustomerID#38497, Country#38498, IsCancelled#38499, Revenue#38500, Year#38501, Month#38502, Day#38503, DayOfWeek#38504, WeekOfYear#38505, Quarter#38506, Hour#38507, IsWeekend#38508, YearMonth#38509], StorageLevel(disk, memory, serialized, 1 replicas)
               +- *(1) ColumnarToRow
                  +- FileScan parquet [InvoiceNo#38491,StockCode#38492,Description#38493,Quantity#38494,InvoiceDate#38495,UnitPrice#38496,CustomerID#38497,Country#38498,IsCancelled#38499,Revenue#38500,Year#38501,Month#38502,Day#38503,DayOfWeek#38504,WeekOfYear#38505,Quarter#38506,Hour#38507,IsWeekend#38508,YearMonth#38509] Batched: true, DataFilters: [], Format: Parquet, Location: InMemoryFileIndex(1 paths)[hdfs://localhost:9000/data/retail/cleaned], PartitionFilters: [], PushedFilters: [], ReadSchema: struct<InvoiceNo:string,StockCode:string,Description:string,Quantity:int,InvoiceDate:timestamp,Un...
```

```
In [37]: import subprocess
# replication on hdfs
subprocess.run(["hdfs", "dfs", "-stat", "%r",
               "/data/retail/cleaned/YearMonth=2010-12/part-00000-4badb10d-eb1b-4871-a5f2-e6f0093826cd.c000.snappy.parquet"], returncode=0)
```

```
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
```

3

```
Out[37]: CompletedProcess(args=['hdfs', 'dfs', '-stat', '%r', '/data/retail/cleaned/YearMonth=2010-12/part-00000-4badb10d-eb1b-4871-a5f2-e6f0093826cd.c000.snappy.parquet'], returncode=0)
```

```
In [ ]:
```